

UNIVERSITY OF CAPE TOWN



EEE4120F

HIGH PERFORMANCE EMBEDDED SYSTEMS

---

## Practicals and Projects

---

27 February 2020

## Introduction

Welcome to the practicals for EEE4120F. These instructions are applicable to all practicals so please take note.

Practicals take place on a Thursday, from 3PM to 6PM. You are responsible for attending these practicals and making use of the resources available. Many of these practicals will require access to specialist hardware that will only be available during the practical session.

It is assumed you have knowledge regarding tools such as git, L<sup>A</sup>T<sub>E</sub>X and running programs from the Linux command line.

It is suggested you have a dual boot system. The reason for this is that development tools are generally better on a Linux based system, whereas proprietary tools are often better supported on Windows. We recommend Ubuntu 18.04 LTS, though swapping between Ubuntu (or any Linux distro) and Windows is an effective strategy, particularly when using proprietary software.

The source code for all pracs is available on the EE-OCW GitHub Project Page: <https://github.com/UCT-EE-OCW/EEE4120F-Pracs>. You can download the source code for pracs from there.

## Prac Overview

Please note that these dates are subject to change. For due dates, refer to Vula.

Table I: The Prac Overview for EEE4120F, 2020

<b>Prac</b>	<b>Dates</b>	<b>Tools</b>	<b>Objective</b>
Prac 1	13/02 - 23/02	Octave	Performance measurement
Prac 2	24/02 - 8/03	OpenCL	SIMD on GPUs
Prac 3	9/03 - 22/03	MPI	Creating image filters on shared mem.
Prac 4	23/03 - 8/04	Vivado	Making a wall time clock
Prac 5	9/04 - 26/04	Vivado	Making a signal generator
Project	27/04-14/05	Vivado	YODA Implementation & demos

# 1 Practical 1 - Testing methodologies

## 1.1 Introduction

The focus of this task is on using OCTAVE (the free sort-of MATLAB program) and doing some statistical operations. In later assignments you will make further use of these statistical functions, and perhaps reuse this code, in comparing and discussing results obtained in other practicals and projects (for example, correlation can be used in analysing gold standard results to higher-speed approximation result).

For installation and tips and tricks on using Octave, visit [The EE Wiki](#) (currently only accessible on the UCT Network - though you can use a VPN).

### Note:

If you're doing a smaller installation of Octave, there are the libraries you require from Octave Forge: audio ; control ; data-smoothing ; fixed ; ga ; gnuplot ; image ; integration ; oct2mat; plot ; signal ; sockets ; specfun ; splines ; statistics ; strings

### 1.1.1 Correlation

Correlation is a useful statistical function for comparing two datasets to judge how similar or different they are. The correlation function returns a correlation coefficient,  $r$ , between -1 and 1. A value of 1 for  $r$  implies perfect positive correlation, i.e. the two datasets are the same. Correlation of 0 implies there is no correlation (the two datasets behave totally differently). A correlation of -1 indicates a total opposite - for example if you compare vectors  $x$  to  $-x$  you get a correlation of -1. Generally if  $|r| \geq 0.8$  there is strong correlation, between 0.5 and 0.8 moderate weak, less than 0.5 is weak (towards no) correlation.

Pearson's correlation (which you can read more about on [Wikipedia](#)) is implemented as follows:

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}} \quad (1)$$

### 1.1.2 Speed Up

$$Speedup = \frac{T_{p1}}{T_{p2}} \quad (2)$$

Where:

$T_{p1}$  = Run-time of original / non-optimal program

$T_{p2}$  = Run-time of optimised program

For obtaining a repeatable timing value, run each version (i.e. initial version and optimised version) of your programs more than once and discard the first measured time. You can, if

you want to be complete, indicate what the initial speed up was and then the average speed up.

### 1.1.3 Critical Section

When measuring performance, it's important to ensure that you are only measuring the section of your algorithm that you need to measure. For example, if you are measuring the execution time of an audio filter, you should not be recording the time taken to open and close the files you are applying this filter to.

## 1.2 Requirements

You are required to run the following experiments:

### 1.2.1 Measuring Execution Time of `rand()`

White noise is often generated with GNU Octave's random number generator `rand()`, which generates uniformly distributed random values in the interval  $[0,1)$ . To create a sound wave of the white noise, the `wavwrite()` function in Octave is used. The `wavwrite()` function expects values in  $[-1.0, 1.0)$ , so the `rand()` output must be multiplied by 2 and shifted down by 1. To generate 10 seconds white noise sampled at 48 kHz, the following instruction are called:

```
white = rand(48000*10,1)*2-1;
```

And to generate a wave file for this noise, we use the `wavwrite()` function as follows:

```
wavwrite(white, 48000, 16, 'white_noise_sound.wav');
```

#### NOTE:

`wavwrite` and `wavread` were deprecated in Octave 5.1. If you're using Octave 5.1 or greater, you need to use `audiowrite` and `audioread`.

```
audiowrite('w.wav', white, 48000, 'BitsPerSample', 16);  
[y, fs] = audioread('w.wav', 16);
```

### 1.2.2 White Noise Generator Script

Next, write a function in a new file called `createwhiten.m` that implements a function with a **for loop** that generates a white noise signal, one sample at a time, comprising N duration in seconds. You can assume that N will always be positive and a multiple of 10. The white noise must be sampled at either 48 kHz or 8 kHz. Name your function `createwhiten(...)`. You

need to use the `rand()` function without arguments so that it will generate a single random value, and the main task is figuring out how to scale so that you create a suitable input to `wavwrite(...)` as explained above. Call the function and check output size as follows:

```
whiten = createwhiten(1000);  
size(whiten);
```

should return:

```
ans = 48000000 1
```

Check that the resulting wave/sound file gives the same sound as the white noise sound.wav generated above by generating the new sound file named white noise sound2.wav and playing it back.

```
wavwrite(whiten, 48000, 16, 'white_noise_sound2.wav');
```

### 1.2.3 Visual Confirmation of Uniform Distribution

Confirm that you've created the sample correctly. Since it's a big signal, let's just look at the first 100 samples by plotting using a histogram function as follows:

```
hist(whiten, 100, 1);
```

This should give image shown in Figure ??:

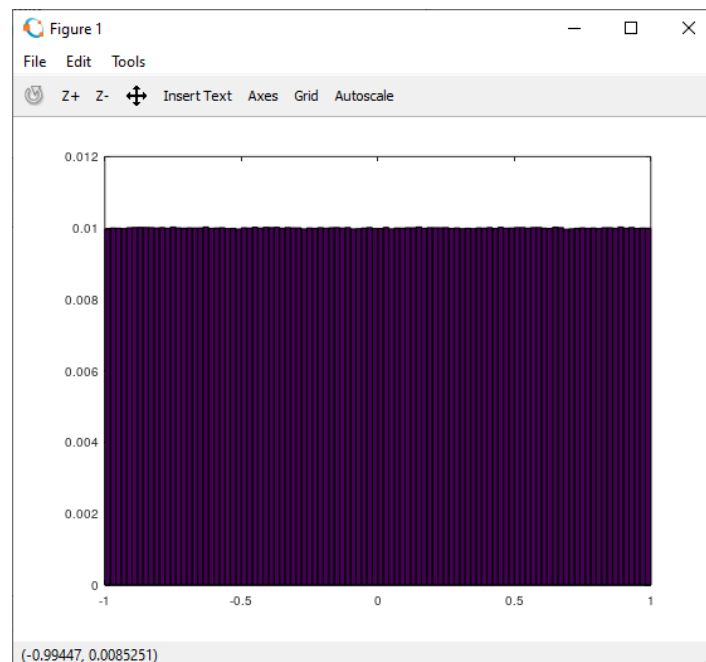


Figure 1: Output of `hist(whiten, 100, 1);`

### 1.2.4 Timing Execution

Now time how long it took to execute the script. Use the `tic()` and `toc()` functions as follows. Note that the call to `tic` is on the same line just before the function - this tends to have less delay between the start of the timer and starting the function. Of course, it is good practice to put it all in a script file.

```
tic; white = rand(48000*1000, 1)*2 - 1; runtime = toc();
disp(strcat("It took: ", num2str(runtime*1000), " ms to run"));
```

Call the `createnwhite(...)` function you created that does the same thing as the `white = rand(48000*1000,1)*2 - 1;` statement that generate the white noise. Measure the time it took for your `createnwhite(...)` function to run and show the timing difference (in milliseconds) and discuss the speed-up you have achieved (if any).

### 1.2.5 Implementing Pearson's Correlation

Implement the Pearson's correlation formula a new m file called it `mycorr.m`. Provide your code in your report. Marks will be awarded on elegance of your code.

### 1.2.6 Comparing Your Correlation Function to Octave

The `cor` function in OCTAVE performs a correlation operation. Compare your `mycorr` results to the build-in `cor` function. First read the noise wave file using `wavread(...)`, then you could, for example, use the following code to test your `mycorr(...)` function as compared to Octave's `cor(...)`:

```
x = wavread('white_noise_sound.wav');
y = x;
r1 = mycorr(x,y)
r2 = cor(x,y) # note that in some versions, this is called "corr"
disp(r2 - r1);
y(1) = 2; y(5) = -4; # i.e. fudge some of the value
r1 = mycorr(x,y)
r2 = cor(x,y)
disp(r2 - r1);
x = rand(1,10); y = rand(1,10);
r1 = mycorr(x,y)
r2 = cor(x,y)
disp(r2 - r1);
```

Generate sample sizes of varying sizes: for example 100, 1000 and 10000 samples. Do a table listing sample sizes vs. `mycorr` speed vs. `cor` speed. Indicate the average speed-up of `cor` to `mycorr`. Include a table of sample size vs time, etc. in your report.

### 1.2.7 Correlation of Shifted Signals

For the last experiment, we want to compare signals shifted in time. Generate sin curves of varying frequency and sampling sizes (again sample sizes 100, 1000 and 10000 samples). Compare samples of the same sizes that are shifted in time, e.g. if A uses `x[1:100]` then B might use `x[11:110]`, in which case B is shifted in time by 10 samples.

For this step only use `cor` to save time. In your report, discuss what you expect the correlation of the identical but shifted signals would be. Run tests to confirm / verify your hypothesis. Provide screen shot plots of some signals you compared.

## 1.3 Submission

Hand in a report about 2-3 pages in length briefly describing your solutions for the tasks above. The format required is the IEEE conference format. Format your report as if it is an article (i.e. don't follow the same chronology as the prac-sheet – format it as “Introduction – Method – Results and Discussion – Conclusion”). In the “Method” section, theorise about what you expect and how you plan on testing said theory. In the “Results” section, confirm that you obtained what you expected (or explain why you obtained something unexpected). Hint: You are trying to answer two questions: 1) “Is my `mycorr` function better suited to run correlation tests, in comparison to the built-in equivalent?” and 2) “Are my theories relating to the correlation of time-shifted sinusoids correct?”



## 1.4 Marking

Table II: Prac 1 Marking Guide

Aspect	Description	Mark Allocation
Report	Intro	2
	Latex/Format	1
	Headings etc	1
	Captions	1
createwhiten	Code intro concepts	4
	Neatness	3
	Structure	3
Sz. v t	Table	2
	Graph	1
	Explanation	2
Shifted signals	Screenshots	2
	Code	2
	Correlation	2
	Explanation	3
	Expected vs Results	1
<b>TOTAL</b>		30

## 2 Prac 2 - OpenCL

### 2.1 Introduction

OpenCL™ (Open Computing Language) is the open, royalty-free standard for cross-platform, parallel programming of diverse processors found in personal computers, servers, mobile devices and embedded platforms. OpenCL greatly improves the speed and responsiveness of a wide spectrum of applications in numerous market categories including gaming and entertainment titles, scientific and medical software, professional creative tools, vision processing, and neural network training and inferencing. [1]

The focus of this practical is investigating the performance metrics of an OpenCL matrix multiplication implementation.

### 2.2 Requirements

The Blue Lab PCs have all the required packages installed, but if you wish to run this practical on your own machine, you will need to install OpenCL as per the instructions on the Wiki - <http://wiki.ee.uct.ac.za/OpenCL>.

### 2.3 The Programming Model

OpenCL uses a programming and memory model similar to OpenGL. The CPU must copy data to the GPU and then tell a kernel (OpenCL worker) to process the data. When the kernel is finished, the CPU can read back the result. Download and inspect the practical source code (available from Vula resources).

Familiarise yourself with the program and what the OpenCL wrapper is actually doing.

If you modify the code, be very careful with how you allocate and free memory. You need to allocate CPU and GPU memory separately, and then free them separately as well.

The local group size has to be an integer fraction of the global group size, in every dimension. Every time you change N, you have to recalculate the local group size.

### 2.4 Tasks

#### 2.4.1 Speed-up

The default program multiplies two 3x3 matrices. This process takes much longer on the GPU than the CPU, mainly due to overhead. At what matrix size does it become worth

while to use the GPU instead of the CPU? Support your argument with measured results (presented in graphs and/or tables). You may also, for instance, make assumptions regarding speed-up expected were the CPU version multi-threaded.

### **2.4.2 Data Transfer Overhead**

The asynchronous nature of the OpenCL interface makes it difficult to obtain accurate timing information of the various steps of the process. Try to come up with a way to measure the data transfer overhead and processing time separately (hint: make use of the `clFinish` function<sup>1</sup> in the OpenCL `WriteData` and OpenCL `ReadData` functions).

Use this new information to comment on the sources of OpenCL processing delay. Also comment on the speed-up factor achieved when transfer overhead is not taken into account. Does this relate well to the number of threads that are running on the GPU? If not, provide an argument for why this is the case. Do this for large  $N$  (pick a value that takes long enough to dominate the transfer overhead, but does not let you finish a whole coffee between runs).

## **2.5 Submission**

Compile your experiments and findings into an IEEE-style conference paper. Make sure you include ALL hardware details, including GPU and CPU clock rates, if available. Of particular importance is the local work group size and number of compute units.

The page limit is 3 pages. Submit your paper to the Vula Assignment for this practical.

## **2.6 Marks**

Note that 33 marks are available, but you will still cannot score a mark above 100%.

---

<sup>1</sup>Which can be found in the `Process_OpenCL()` function in the `OpenCL_Wrapper.c`

Table III: Prac 2 Marking Guide

Aspect	Description	Mark Allocation
P1	Graphs/Tables *	6
	Comparison to CPU Speed up	3
P2	3 Metrics reported on	4
	Use of clFinish in Code	2
	Overhead discussion	3
	Discussion	4
General	Introduction	3
	Layout/Captions etc	3
	PC details	2
Bonus experiments		3
TOTAL		30

## 3 Prac 3 - MPI

### 3.1 Overview

The focus of this practical is message-passing in a distributed memory programming model. The processing algorithm is a median filter will be used. You can find a tutorial at <https://computing.llnl.gov/tutorials/mpi/>, but this is much more detailed than what you need for this practical.

### 3.2 The Programming Model

One thing to keep in mind is that each process runs in a completely separate memory-space. There are no shared memory structures, so any global memory you define is local to each instance. The only means of sharing data is through sending and receiving messages. Typically MPI is only useful on distributed systems (systems where each node has its own memory, and the nodes are connected through some form of network – typically Ethernet). Setting up such a system is not trivial, and therefore not included in this practical. For this practical, all the processes will be running on the same CPU, with memory being separated by means of memory virtualisation (through use of the [memory management unit](#)).

#### 3.2.1 MPI Commands

MPI programs are compiled with **mpic**, not gcc, and run with **mpirun**, not by calling the program binary directly. The makefile associated with this practical runs the program with 5 instances (one master and 4 slaves). Feel free to change this if you so wish.

Study the source code provided. It has been adapted from [this example](#) and uses simple blocking communication. You can implement this practical by means of blocking communication, so you don't need any more commands than those in the example code.

If you feel like being fancy, have a look at the [MPI Iprobe](#) command.

You can also have a look at [this link](#) to see how to tag your messages. Tags are typically used to determine the message type before actually receiving it, but it could be used for many other things as well.

There are many MPI datatypes, but for the purpose of this practical you can get away with using bytes (MPI BYTE) exclusively.

### 3.3 Problems

#### 3.3.1 Problem Partitioning

Choose a memory partitioning scheme. In your report, explain your partitioning scheme and why you chose it. Do not send the entire dataset to each process, as this is a waste of data. **[5 marks for describing the partitioning]**

You also need to have some means to indicate to the slave what size the data is. You can implement this any way you wish, but it is suggested that you send one message, with a known size and fixed format, that indicates data size and any other parameters you wish to send to the slave. The next message can then be the actual data to work on. **[5 marks for getting data to slaves]**

When the slave is finished, it must send the result back to the master. The master must then assemble the results and save the resulting filtered image to disk. **[5 marks sending back the data and reassembling]**

Provide snippets of your code in the report and attached your code in a zip repository (it is important for your code snippets to be suitable commented, the use of comments in your main code repository is not as important for marking but is good practice). **[5 marks for code]**

#### 3.3.2 Experimentation

Feel free to experiment with different number of processes, although this is not required. It is sufficient to simply get it working. **[5 marks for experiments]**

Your report should comment on the time performance though, so make sure you take time measurements. **[5 marks graphs and display of results]**

### 3.4 Submission

Compile your experiments and findings into an IEEE-style conference paper. The page limit is 3 pages. Submit your paper to the Vula Assignment for this practical.

### 3.5 Marking

Note that 33 marks are available, but you cannot score more than a total of 100%.

Table IV: Prac 3 Marking Guide

Aspect	Description	Mark Allocation
Partitioning	Listing (show data split)	2
	Explanation	3
To Slaves	Listing & ACK	2
	Explanation	3
On/From Slaves	Listing	3
Reassembling	Explanation	2
Code	Comments	2
	Compiles & runs	3
Experiments	Experiment	3
	Good practice	2
Results	Graph/Table	2
	Discussion	3
Bonus		3
<b>Total</b>		<b>30</b>

## 4 Prac 4 - FPGA Introduction

### 4.1 Introduction

For this practical, it is important to complete the tutorial in order to get acquainted with the tools required to complete the practical. It is recommended that you work in groups of 3, and that all three members work through the tutorial individually.

It is also imperative that you read through the wiki. Operation of Vivado can be intimidating at first use. The Wiki has been carefully written to include all you need for the practicals.

### 4.2 Tutorial

The tutorial is not for marks but it is suggested you complete it, as it will give you the basics of Vivado, and allow you to upload a simple program to the FPGA. In the tutorial (as on the wiki) we use the Digilent Nexys A7 as an example board, but the process should be the same for the Nexys 4 DDR and Nexys 4. The only thing that will change between the three is the board you select when creating the project, and the constraints file you use. Note the constraints file does determine naming of some of the I/O, so double check that.

Most of the details on how to complete these steps are on the wiki under [Xilinx:Vivado](#).

1. Install Vivado
2. Install boards

The only boards you might work with in this course are the Nexys 4, Nexys 4 DDR, and the Nexys A7. So you only need to worry about installing those.
3. Download and save the constraint files. That, or copy the simple example given on the wiki if you're using the A7.
4. Create a new project.
  - You can call it "Tutorial".
  - Set it as an RTL project and select "Do not add sources at this time". For the prac later, you can add the given source files here.
  - Select your board appropriately.
5. Add constraint source
  - Right click on constraints, select "Add sources". In the dialog box, make sure "Add or create constraints" is selected. Hit next.
  - Select "Create file" if you are copy pasting the constraints in, or "add files" if you have the constraints file downloaded locally. Press "Finish".



- If you were intending on copy-pasting constraints in, do so now. Ensure your constraints are correct for the board, and have the clock, two switches and two LEDs enabled.
  - Right click on the constraints file, and select "Set as Target Constraint File"
6. Add the Verilog source
    - Right click on "Design Sources, select "Add sources". In the dialog box, make sure "add or create constraints" is selected. Hit next.
    - Select "Create file". Call it the same name as your project (good practice for the top level module). Press "Finish".
    - A dialog will open, with ports. You can just press "okay", as we'll define ports in the Verilog code.
    - Right click on the Verilog file, and select "Set as Top" (if it is not already set as top - indicated by being in bold).
    - In it, paste code to, each clock cycle, write switch[0] to LED[0] and the inverse of switch[1] to LED[1]. Example code can be found on the wiki.
  7. Select "Run Synthesis"
  8. Select "Run Implementation"
  9. Select "Generate Bitstream"
  10. Upload your bitstream to your target board (speak to a tutor to get a board)
    - (a) Plug in the board
    - (b) Select "Open Hardware Manager"
    - (c) Select "Open Target" and then "Auto Connect"
    - (d) Vivado should find the board. Select "Program Device" and select the board you've plugged in (The A7 shows as "xc7a100t\_0". Press the "Program" button.
  11. Hooray! You've created your first FPGA circuit. Toggle the switches to see if it operates as expected. Now let's move on to the fun stuff.

## 4.3 Practical

### 4.3.1 Introduction

In this practical, you will create a digital clock on an FPGA. There is no report for the practical, but you will need to submit screenshots of your testbenches and demonstrate your implementation to a tutor.

Source files are available on the EEE4120F OCW GitHub: <https://github.com/UCT-EE-OCW/EEE4120F-Pracs>.

### 4.3.2 Given Modules

1. TLM  
The top level module, called "Clock.v" in the source files on GitHub, contains the primary logic for your wall clock and allows you to implement I/O and other modules
2. Delay\_Reset  
It's also useful as many components require a set up time. So by using a delayed reset signal, we can cater for reset times of peripherals.
3. Seven-Segment Driver  
This module takes 4 BCD values and displays them on the seven segment display.
4. Decoder  
Used by the Seven-Segment Driver to decode decimal to the appropriate cathode pins.
5. Debounce  
A debounce module you'll need to implement in order to debounce button presses.
6. PWM  
A module you'll need to implement in order to give the seven segment displays changing brightness. This can be tricky, it's suggested you leave it for last.

### 4.3.3 Requirements

The following outcomes are required to pass the demonstration:

1. Implement a simple state-machine to display the real time (hours and minutes) on the 7-segments display. You can start your clock at 00:00 upon reset. Make your clock faster in order to test that the time overflows correctly. Use a 24-hour time format.  
  
The easiest way to do this is with deeply nested if statements. Run a counter that overflows every second and increment the seconds counter on every overflow. Every time this seconds counter equals 59 (i.e. it will overflow on this clock-cycle), increment a minutes units counter. Every time this minutes units counter is about to overflow, increment a minutes tens counter, etc.  
  
Only use non-blocking assignments ( $\leq$ ). Blocking assignments ( $=$ ) inside clocked structures are much more difficult to debug. Remember that the entire always block is evaluated at once: the statements are not evaluated sequentially.
2. Display the seconds on the LEDs, in binary format. This is done with a simple assignment outside the always block.
3. Use one of the buttons, properly debounced, to set the minutes. It must increment time by one minute every time it is pressed. Make sure that your time overflows correctly. You do not need to increment the hours when changing the minutes.  
  
It is recommended to write a Debounce module for this. On every clock cycle of the system clock (the fast one), check the state of the button. If it is not the same as the

current module output, change the output and start a dead time counter. While this counter is counting, do not change the output of the module, no matter what the input is doing. Use a deadtime of between 20 ms and 40 ms. To prevent unstable states, register the button before use.

In the clock state machine, you can use a register to store the button's 'previous' state. If the current state is high, and the previous state is low, the button signal went through a rising-edge. Do not use always @(posedge Button) – keep everything in the same clock domain.

4. Use another one of the buttons, properly debounced, to set the hours. It must increment time by one hour every time it is pressed. Make sure that your time overflows correctly.
5. Use the slide-switches to represent a binary "brightness" word. Make use of pulse-width modulation (PWM) to dim the brightness of the LED display.

Ensure that the phasing between the driver signals and the PWM signals is correct. The easiest way to do this is to select a PWM frequency such that the PWM signal goes through exactly one period between driver signal state changes. You can implement this within the SS Driver module.

#### 4.4 Mark Allocations

Table V: Prac 4 mark allocation

		<b>Marks</b>
<b>Testbench</b>		12
	Minute seconds reaching 60 and increasing minutes	
	Minutes reaching 59 and increasing hours	
	Time reaching 23:59 and wrapping back to 00:00	
	(3 each +3 for neatness)	
<b>Demo</b>		
	Time flows correctly (minutes overflow to an increase in hours, and 23:59 flows to 00:00) [3, subtracting 1 for each missed objective]	3
	Time can be scaled through a variable (i.e. the count to increase seconds isn't fixed)	1
	Minute button increases minutes and doesn't increase hours	1
	Hour button increases hours	1
	Debounce module implemented	2
	PWM module	5
	<b>TOTAL</b>	25

**For the 5 marks on PWM:** 1 mark for attempted, 2 marks for reading switches and adjusting, 3 marks for "flashing" implementation, 4 marks for some mix between flashing and decent PWM, 5 marks for correctly implemented.

## References

- [1] “Opencl - the open standard for parallel programming of heterogeneous systems,” Jul 2013. [Online]. Available: <https://www.khronos.org/opencl/>